





BUILD

//LOCALHOST:LONDON

Wednesday, June 24, 2026 · 08:45–13:30
Storey Club, Paddington






Who's running your agents?

Governing AI agents as first-class identities with Microsoft Agent 365

Mika Vilpo
Azure Lead Architect, Netox Oy
Microsoft MVP - Cloud Security



MIKA VILPO

IT GEEK BUILDING CLOUDS

@NETOX OY

LIKES TO TALK ABOUT IDENTITY, AI, AUTOMATION AND
SERVERLESS

25+ YEARS IN IT AND 15+ YEARS IN CLOUD

SPENDING FREE TIME AS RED CROSS VOLUNTEER IN
DISASTER MANAGEMENT, FIRST AID AND SEARCH & RESCUE

MIKA.VILPO@NETOX.COM

@MIKAVILPO 



Every agent is a non-human identity you do not yet govern

Agents authenticate and call resources like any workload, but they are created ad hoc across Copilot Studio, Foundry, open-source frameworks and local endpoints.

INVISIBLE

No inventory of agent identities or their consented grants

OVER-PERMISSIONED

Broad application permissions, standing consent, no expiry

UNATTRIBUTABLE

Shared app principals hide which agent did what

Shadow IT, round two. This time the principals act on their own.

1.3B


AI agents projected by 2028

each one a non-human identity to govern

IDC · via Microsoft

The compliance clock is already running

EU AI ACT · PHASED ENFORCEMENT



REGION	APPROACH	WHERE IT STANDS
EU	Risk-based, binding	Agents in scope via GPAI; high-risk obligations from Aug 2, 2026
UK	Principles-based, sectoral	No single AI Act; CMA confirmed existing law already applies (Mar 2026)

Where Agent 365 fits: the operational controls these regimes expect (inventory, traceability, oversight, audit). An enabler, not compliance in a box.

/ WHAT IT IS

Agent 365 is the control plane for agents

One place to discover, govern and secure every agent, regardless of where it was built or hosted.



OBSERVE

Registry, usage and risk for every agent before it bites.



GOVERN

Lifecycle, access policy and audit, enforced not just reported.



SECURE

Agent identities, data protection and runtime threat defence.

Surfaced through M365 admin center, Entra, Purview and Defender, the tools you already operate.

Five capabilities



Registry

Inventory of every agent



Access control

Risk-based access via Entra



Visualization

Map agents, users, resources



Interoperability

MCP and A2A across the estate



Security

Identity, DLP, threat defence

Framing introduced at Ignite 2025 by Charles Lamanna: registration, access control, visualization, interoperability, security.

Built on the tools you already run

AGENT 365 control plane · observe / govern / secure

↑ enforcement points ↑

M365 admin center

registry · Agent Map

Entra Agent ID

identity · CA · network

Microsoft Purview

DLP · labels · audit

Microsoft Defender



Cloud · Cloud Apps · XDR

YOUR AGENTS Microsoft · open-source · third-party · multicloud

Each layer is a real enforcement point, not a parallel system. You extend the controls you already run for people.




Demo





Agent Identities




Where each agent's identity comes from

Entra Agent ID is the common primitive, but how an agent gets one depends on where it was built.

BUILD SURFACE	PERSONA	ENTRA AGENT ID
Agent Builder (M365 Copilot)	no-code, end users	none today
Copilot Studio	low-code, makers	opt-out per environment
Microsoft Foundry	pro-code, enterprise	automatic
M365 Agent SDK	pro-code, any runtime	provisioned by the SDK
Third-party / multicloud	external / synced	manual or via SDK

Build platforms stay separate; Agent 365 governs the result. Foundry auto-provisions, the SDK provisions for any framework.

The Entra Agent ID object model



TWO ARCHETYPES

Assistive / interactive

Delegated permissions, user token (scp claim).
Acts on behalf of the signed-in user.

Autonomous

Application permissions, app-only token (roles claim).
Acts under its own identity, in the background.

Credentials live on the blueprint; agent identities hold none and are impersonated by it.

When to create an agent blueprint

Default rule: one blueprint per trust boundary (an agent type or family). One agent identity per instance.

Blueprint granularity

New blueprint when:

- Different trust boundary (separate runtime, secrets, file system)
- Different agent type or publisher
- Different shared Conditional Access or permission posture

Not a reason for a new blueprint:

- Disabling one agent (target it with Conditional Access)
- Separate audit (each identity already logs on its own)
- Replicas of the same agent (that is scaling, not identity)

One identity per agent instance, up to 250 per blueprint.

Where permissions live

On the blueprint (shared, inherited by all identities):

- Delegated scopes, app roles, API permissions
- Conditional Access, authentication and credentials

On the agent identity (per instance):

- Extra access via access packages: app permissions, group memberships, Entra roles

Blueprint grants A, B, C -> every identity inherits A, B, C. An identity can also be granted **D** directly via access packages. The blueprint is the floor, not the ceiling; to grant less, use a narrower blueprint.

How this maps to what you already know

Two familiar objects, one new one (the agent), and where the managed identity fits.

The object mapping

App registration -> **Agent identity blueprint**
the template and definition

Enterprise app (SP) -> **Blueprint principal**
live tenant instance: consent, sign-in, provisions

(no classic match) -> **Agent identity**
the running agent: per instance, holds its own RBAC

User account -> **Agent user**
optional, strict 1:1

Where the credential lives

Managed identity (user-assigned)

federated identity credential (FIC)



Blueprint principal

authenticates to Entra, no stored secrets



Agent identity

acts on resources, holds Azure RBAC




Azure resource

RBAC evaluates the agent identity, not the MI

The managed identity authenticates the blueprint, not the resource.

Worked example: a Foundry agent acting in Azure



InvoiceOps reads invoice blobs and posts to an ERP through an MCP server. Watch where the managed identity stops.



The managed identity stops at the blueprint. RBAC and audiences live on the agent identity.



Demo



/ FOR DEVELOPERS

Make any agent Agent 365-native

Framework-agnostic SDK plus CLI. Wrap your agent, create its blueprint, inherit governance.



Identity

Entra Agent ID + mailbox



Observability

OpenTelemetry spans



Notifications

Teams, Outlook, email



Governed access

Governed MCP to M365 data

```
install
# python
pip install microsoft-agents-a365-runtime
# node
npm i @microsoft/agents-a365-runtime
# .net
dotnet add package Microsoft.Agents.A365.Runtime
```

WORKS WITH ANY STACK

Agent Framework

M365 Agents SDK

OpenAI Agents SDK

LangChain

Claude Code SDK

Azure AI Foundry

Registration = wrap with the SDK, then the Agent 365 CLI creates the Entra App Blueprint with the MCP servers it needs.

What's free vs what you pay for

CAPABILITY	FREE	A365 LICENCE
Agent identity (Entra Agent ID) + blueprints	✓	✓
Agent inventory and visibility	✓	✓
Basic usage insights	✓	✓
Core admin governance actions	✓	✓
Advanced usage analytics	X	✓
Agent governance policies	X	✓
Security posture management	X	✓
Threat detection	X	✓
Data security controls	X	✓

Who needs a licence, and how to buy

SKU	PRICE	INCLUDES
Agent 365 (standalone)	\$15 /user/mo	The agent governance control plane
Microsoft 365 E7	\$99 /user/mo	E5 + Copilot + Agent 365 + Entra Suite

Both generally available since 1 May 2026. Licence is per human who interacts with, manages or sponsors an agent.

WHICH	Broad Copilot rollout → E7. Small control-plane team → standalone.
PREVIEW	CA + ID Protection for agents: GA for delegated (on-behalf-of), public preview for own-identity.

The per-seat price is the governance floor. Building and running agents is separate Copilot Studio / Foundry consumption (Copilot Credits).

Previews and roadmap

CAPABILITY	STATUS
Local agent discovery (Defender + Intune)	PUBLIC PREVIEW
Policy controls + runtime blocking	PUBLIC PREVIEW
Purview runtime DLP for agent prompts (Foundry)	PUBLIC PREVIEW
Own-identity Conditional Access + ID Protection	PUBLIC PREVIEW
Multicloud sync: Bedrock, Vertex AI, Agentforce, Genie	PUBLIC PREVIEW
Autonomous agentic-user identities	FRONTIER · no GA date
Microsoft Scout (first Autopilot agent)	FRONTIER · GA early 2027
Agent 365 SDK	GA · Build 2026

1 Jul 2026: AIAgentsInfo → AgentsInfo; RTP rules move to the new Policies experience (Block stops unless redefined); single Security for AI Agents toggle.

What to do on Monday

01

See what you already have

The registry populates itself. Start there.

02

Design blueprints, scope permissions

Inheritable perms and sponsor accountability before deploy.

03

Mind the delegated vs own-identity line

Own-identity CA and ID Protection are still preview.

04

Budget beyond the seat

Per-seat is governance; run cost is Copilot Credits.

Thank you

Questions?

Agent 365 overview	learn.microsoft.com/microsoft-agent-365
Entra Agent ID (identity model)	learn.microsoft.com/entra/agent-id
Developer docs (SDK + CLI)	learn.microsoft.com/microsoft-agent-365/developer
Build BRK251 / OD840	build.microsoft.com
Plans & pricing	microsoft.com/microsoft-agent-365

Mika Vilpo, @mikavilpo